


# **CIT650 Introduction to Big Data**

## **Introduction to SPARK**

# Distributed Computing vs. Parallel computing


## Parallel Computing

processors access shared memory



## Distributed Computing

processors usually have their own private memory



## Distributed Computing Benefits

- Distributed systems are inherently scalable as they work across multiple independent machines and scale horizontally.
- Distributed systems offer fault tolerance as independent nodes fail without affecting system integrity.
- Distributed systems provide redundancy that enables business continuity.

# Spark for Distributed Computing

- Spark supports a computing framework for large-scale data processing and analysis.
- Spark provides parallel distributed data processing capabilities.
- Spark provides scalability.
- Spark provides fault-tolerance on commodity hardware.
- Spark enables in-memory processing.
- Spark enables programming flexibility with easy-to-use Python, Scala, and Java APIs.

# Spark vs. MapReduce



# Apache Spark<sup>1</sup>

- A general-purpose, fast, large-scale data processing engine
- Started in AMPLab @UC Berkley, now Data Bricks
- Written in Scala
- Considered as a third-generation distributed data processing system
  - Hadoop is considered as a second generation
- Why would we need a new generation?
  - Limitations of MapReduce (Hadoop)
    - Materialized intermediate results
    - Abstraction of Map and Reduce Applications
    - Can only be written in Java
    - Poor support for real-time data processing
  - Exploit advancements in hardware
    - Memory is much cheaper
    - Multi-core is now a commodity



10x faster than Hadoop on disk  
100x faster than Hadoop in memory

2-5x less code


<sup>1</sup>Zaharia, Matei, et al. "Spark: Cluster computing with working sets

# Third Generation Distributed Systems

- Handle both batch and real-time processing
- Exploit RAM as much as disk
- Multiple-core aware
- Do not reinvent the wheel
  - Use HDFS for storage
  - Apache Mesos/YARN for execution
- Plays well with Hadoop
- Iterative processing

# Functional Programming

- Mathematical **function** programming style
- Follows a declarative programming model
- Emphatizes **what** instead of **how to**
- Uses **expressions** instead of **statements**





# Functional Programming

**Scala code** defines a lambda expression with two parameters x and y of type Int. The lambda expression is stored in the variable add.

## Scala

```
// an example lambda  
operator in Scala to add  
2 numbers  
val add = (x:Int, y:Int)  
=> x + y  
println(add(1,2))
```

**Python code** defines a lambda function with two parameters x and y. The lambda function is assigned to the variable add.

## Python

```
// an example lambda  
function in Python to add  
2 numbers  
add = lambda x, y : x + y  
print(add(1,2))
```

# Functional Programming – Passing Operations

## Scala

```
// Define the function
def performOperation(x: Int, y: Int, operation: (Int, Int) => Int): Int = { operation(x, y) }

// Call the function with a lambda
val result = performOperation(5, 3, (a, b) => a * b)
println(result) // Output: 15
```

## Python

```
# Define the function
def perform_operation(x, y, operation): return operation(x, y)

# Call the function with a lambda
result = perform_operation(5, 3, lambda a, b: a * b)
print(result) # Output: 15
```

# Functional Programming – using Arrays


## Scala

```
val add = (x: Int, y: Int) => x + y
val numbers = List((1,2), (3,4), (5,6))
val sums = numbers.map (pair => add(pair._1, pair._2))
println(sums) // Prints: List(3, 7, 11)
```

## Python

```
add = lambda x, y: x + y
numbers = [(1,2), (3,4), (5,6)]
sums = list(map(lambda pair: add(pair[0], pair[1]), numbers))
print(sums) # Prints: [3, 7, 11]
```

# Unified Platform for Big Data Processing



# Why Unification?

- Good for developers: One platform to learn
- Good for users: Take apps everywhere
- Good for distributions: More applications
- Is based on a common abstraction

# Spark Abstractions

- Spark core abstraction is Resilient Distributed Dataset (RDD)
  - Resilient: fault tolerant and can be recomputed when recovering from a failure
  - Distributed: processing takes place over several nodes in parallel, like MapReduce
  - Dataset: initial data can come from files, memory, or created programmatically
  - Immutable: once created cannot be changed
  - Lineage: each RDD knows about its parents
- Spark applications are series of operations that transform input RDDs into output RDDs or final values

# Word count In Spark

- Recall how we defined the code for word count in MapReduce?
- How does it look in Spark?

```
sc = SparkContext(appName="PythonWordCount")

lines = sc.textFile(sys.argv[1])



counts = (lines.flatMap(lambda x: x.split(' ')) # Split lines into words
          .map(lambda x: (x, 1))                # Map each word to a tuple (word, 1)
          .reduceByKey(lambda a, b: a+b)) # Reduce by key, sum occurrences

output = counts.collect()

for (word, count) in output:
    print("%s: %i" % (word, count))
sc.stop()
```

# RDD Operations

- Two main types of RDD operations
  - Transformations: result in a new RDD
    - Can be chained
    - Forked
    - joined
  - Actions: return values, no more RDDs
    - One action at the end of each transformation chain





# RDD Transformations

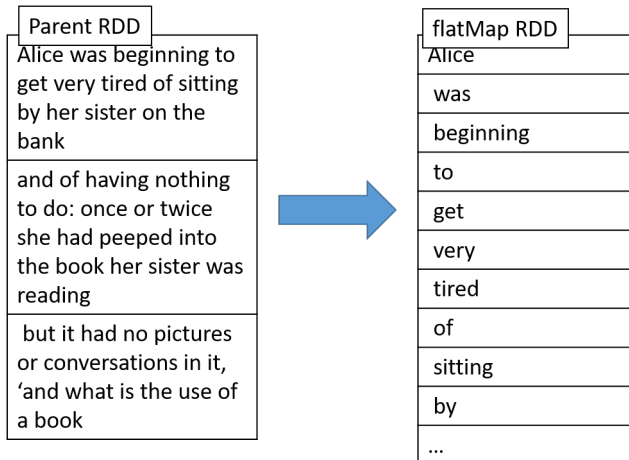
- Transformations create a new RDD from an existing one
- RDDs are immutable
  - Apply a series of transformations to modify the data as needed
- Common transformations
  - Map(function): 1-to-1 mapping
  - FlatMap(function): 1-to-Many mapping
  - Filter(function): 1-to-1 mapping with selectivity
- Special transformations
  - ReduceByKey
  - GroupByKey

## Map vs. FlatMap

Aspect	Map	FlatMap
Functionality	Transforms each element of a collection independently.	Transforms each element and then flattens the result.
Input Example	"Hello world", "This is a test"	"Hello world", "This is a test"
Operation	Splits each line into words.	Splits each line into words and merges them into a single collection.
Output	[[ "Hello", "world"], [ "This", "is", "a", "test"]]	[ "Hello", "world", "This", "is", "a", "test"]
Structure	Collection of collections (array of arrays).	Single flat collection (array).

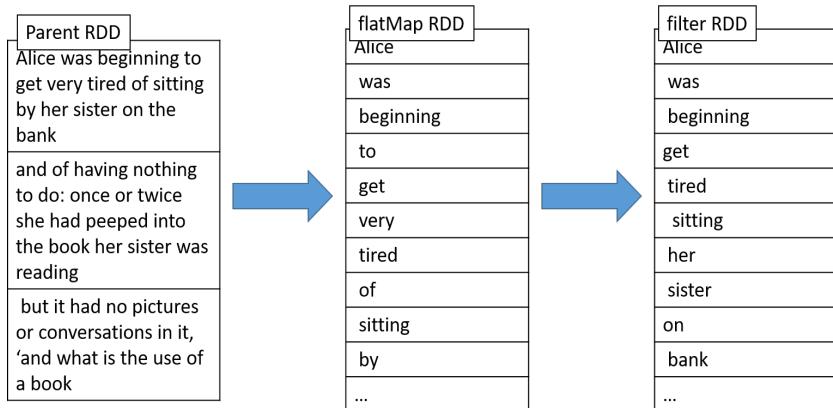
## Transformation: flatMap

```
lines = lines.flatMap (lambda x : x.split (' '))
```



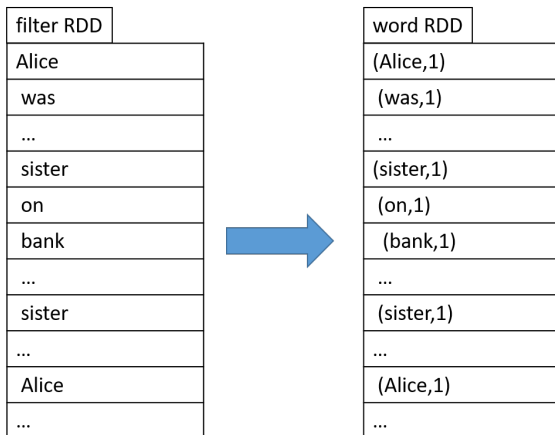
## Transformation: filter

```
filtered = lines.filter (lambda x : x not in ['by','very','to',  
'the'])
```



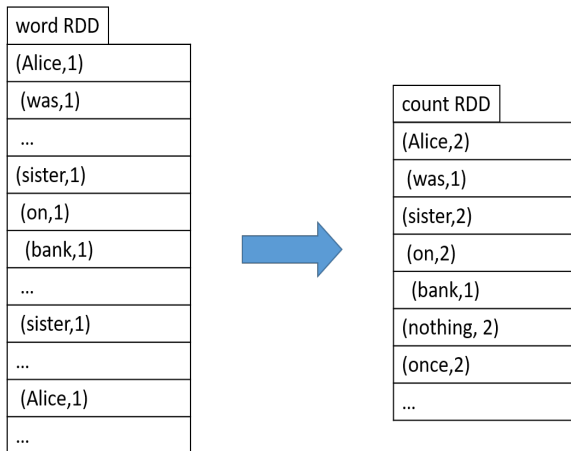
## Transformation: map

```
word = filtered.map(lambda x : (x,1))
```



## Transformation: reduceByKey

```
counts = word.reduceByKey(lambda x,y : x + y)
```



# RDD Actions

- Actions trigger execution of transformation chains
- No further RDD transformations
- Common actions
  - `Collect()`: returns an array of all the elements
  - `Take(n)`: returns an array of the first n elements
  - `Count()`: returns the number of elements in RDD
  - `saveAsTextFile()`: saves the data to file system, either HDFS for local

## Action: collect

```
output = counts.collect( )
```

count RDD
(Alice,2)
(was,1)
(sister,2)
(on,2)
(bank,1)
(nothing, 2)
(once,2)
...



```
output=  
[(Alice,2), (was,1), (sister,2), (on,2), (bank,1), (nothing,2), (once,2),...]
```



# Lazy Execution

- Spark follows a lazy execution scheme
  - No RDDs are computed until an action is specified
- Why?
  - Help optimize execution plan
- Only lineage is created as moving from one transformation to the other
  - To learn about lineage of a chain of transformations, call `toDebugString()` after the transformation you are interested in

# RDD Lineage


## Application

```
mydata_filt = sc.textFile('file.txt')  
.map(lambda line: line.upper())  
.filter(lambda line: line.startswith('I'))
```

## Lineage

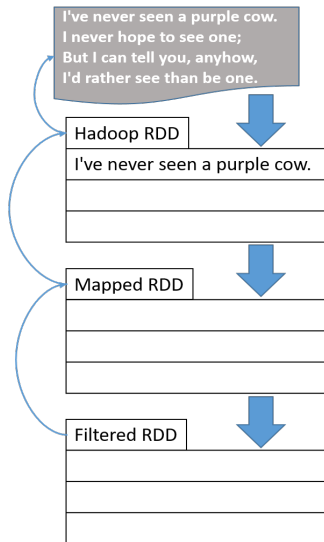
```
mydata_filt.toDebugString()
```

```
(2) FilteredRDD[7] at filter ...  
| MappedRDD[6] at map ...  
| file.txt MappedRDD[5] ...  
| file.txt HadoopRDD[4] ...
```



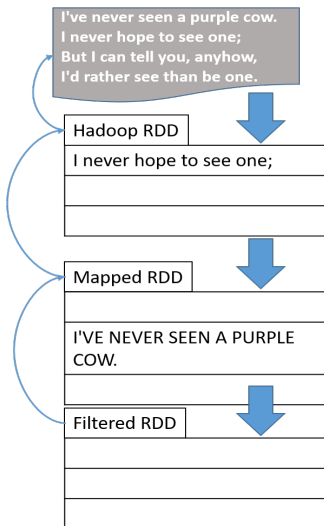
# Pipelining

- When possible, Spark will pass individual outputs of each transformation to the next.
  - In Hadoop, all intermediate results are completely calculated before beginning the next step
- Pipelining helps reduce latency



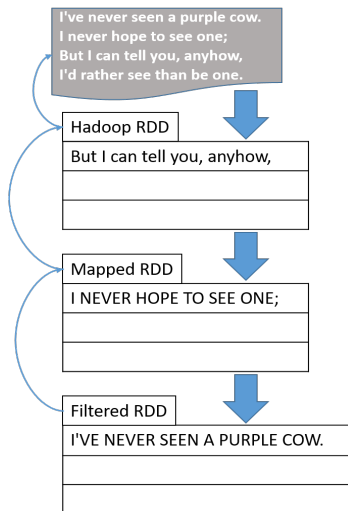
# Pipelining

- When possible, Spark will pass individual outputs of each transformation to the next.
  - In Hadoop, all intermediate results are completely calculated before beginning the next step
- Pipelining helps reduce latency



# Pipelining

- When possible, Spark will pass individual outputs of each transformation to the next.
  - In Hadoop, all intermediate results are completely calculated before beginning the next step
- Pipelining helps reduce latency



# Creating RDDS

- We learned that RDDs can be created as a result of transformations on parent RDDs
- What about root RDDs?
- Can be created from
  - Data in memory, collections
  - From files, example from text files as we saw earlier

# Creating RDDs from Collections

- `SparkContext.parallelize(collection)`

```
mydata = ['Alice' , 'Jack' , 'Andrew' , 'Frank']  
myRDD = sc.Parallelize(mydata)  
myRDD.take(2)  
output : ['Alice' , 'Jack']
```

- Useful for:
  - Testing
  - Integrating

# Creating RDDs from Files

- So far, we saw `sc.textFile("file")`
  - Accepts a single file, a wildcard list of files, or comma-separated list of file names
  - Examples:
    - `sc.textFile("myfile.txt")`
    - `sc.textFile("mydata/*.log")`
    - `sc.textFile("myfile1.txt,myfile2.txt")`
    - `textFile` only works with line-delimited text files
  - Each line in the file is a separate record in the RDD
  - Files are referenced by relative or absolute URI
    - Absolute URI: `file:/home/training/myfile.txt` or `hdfs://localhost/loudacre/myfile.txt`
    - Relative URI (uses default file system): `myfile.txt`
  - What about other file formats?



# Creating RDDs from Other File Formats

- Spark uses Hadoop's InputFormat and OutputFormat Java classes
  - TextInputFormat/TextOutputFormat
  - SequenceInputFormat/SequenceOutputFormat
  - FixedLengthInputFormat
- Support for other formats
  - AvroInputFormat/AvroOutputFormat

# Using Input/output Formats

- Define input format using `sc.hadoopFile`
  - Or `newAPIHadoopFile` for New API classes
- Define output format using `rdd.saveAsHadoopFile`
  - Or `saveAsNewAPIHadoopFile` for New API classes

## Example:

```
input_rdd = sc.newAPIHadoopFile(  
    "path/to/textfile.txt",  
    "org.apache.hadoop.mapreduce.lib.input.TextInputFormat",  
    "org.apache.hadoop.io.LongWritable",  
    "org.apache.hadoop.io.Text" )
```

# Whole-file-based RDDs

- `sc.textFile` puts each line as a separate element
  - What if you are processing XML or JSON files?
- `sc.wholeTextFile(directory)` creates a single element in RDD for the whole content of a file in the input directory
  - Creates a special type of RDDs, (paired RDDs), we discuss later
  - Works for files with small sizes (elements must fit in memory)

file1.json

```
{ "id": "123",  
  "name": "Ahmed",  
  "score": 717  
}
```

file2.json

```
{ "id": "312",  
  "name": "Mark",  
  "score": 810  
}
```

:

Whole RDD
(file1.json, { "id": "123", "name": "Ahmed", "score": 717 })
(file2.json, { "id": "312", "name": "Mark", "score": 810 })
...

# RDD Content

- RDD can hold elements of any type:
  - Primitive data types
  - Sequence types
  - Scala/Java objects (if serializable)
  - Mixed types
- Special RDDs
  - Pair RDDs: consist of key-value pairs,
    - recall map step of word count example
    - `sc.wholeTextFile`
  - Double RDDs
    - RDDs consisting of numeric data

# Other General RDD Transformations

- Single RDD transformations
  - distinct: removes duplicate values
  - sortBy: sorts by the input function
- Multi-RDD transformations
  - intersection: outputs common elements of the input RDDs
  - union: add all elements from input RDDs to the output RDD
  - zip: performs a cross product between the input RDDs

[1, 2, 3] zip ['a', 'b', 'c']  [(1, 'a'), (2, 'b'), (3, 'c')]

# Pair RDDs

- A special form of RDDs
  - Elements must be a tuple of two elements (key, value)
  - Keys and values can be of any type
- Why use Pair RDDs
  - To have the benefits of MapReduce
  - Ability to scale by forwarding tuples of the same key value to the same processing node, shuffling.
  - Other additional transformations are built-in, e.g., sorting, joining, grouping, counting, etc.

# Creating Pair RDDs

- You can have your root RDD as a pair RDD, e.g., `sc.wholeTextFile`
- You can use a transformation to put the data in pair RDDs
  - `map`
  - `flatMap/flatMapValues`
  - `keyBy`

## Example

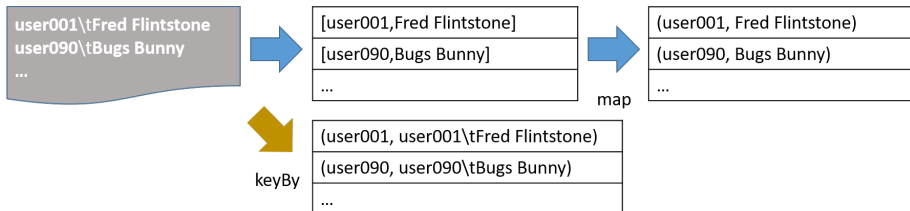
- Create a pair RDD from a tab-delimited file

```
users=sc.textFile (file)
```

```
.map(lambda line: line.split("\t"))
```

```
.map(lambda elems: (elems[0],elems[1]))
```

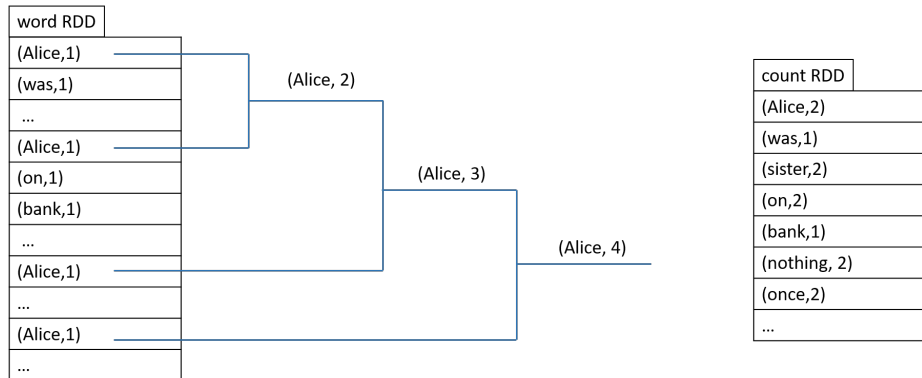
```
.keyBy(lambda line : line.split('\t')[0])
```





# Transformation: reduceByKey

```
counts = word.reduceByKey( lambda x, y: x+y )
```



## Other Pair RDD Transformations


- `countByKey`
  - Returns a pair RDD with the same key as parent and value is the count of key occurrences
- `groupByKey`
  - Similar to the input of Hadoop Reducer, (key, [list of values])
- `sortByKey(ascending=true/false)`
  - Returns a pair RDD sorted by the key
- `join`
  - Takes two input pair RDDs with the same key (key, value1), (key, value2)
  - Returns (Key, (value1,value2))

# Examples

(ord001, sku101)
(ord001, sku190)
(ord001, sku030)
(ord002, sku101)
(ord002, sku912)
(ord003, sku999)
...



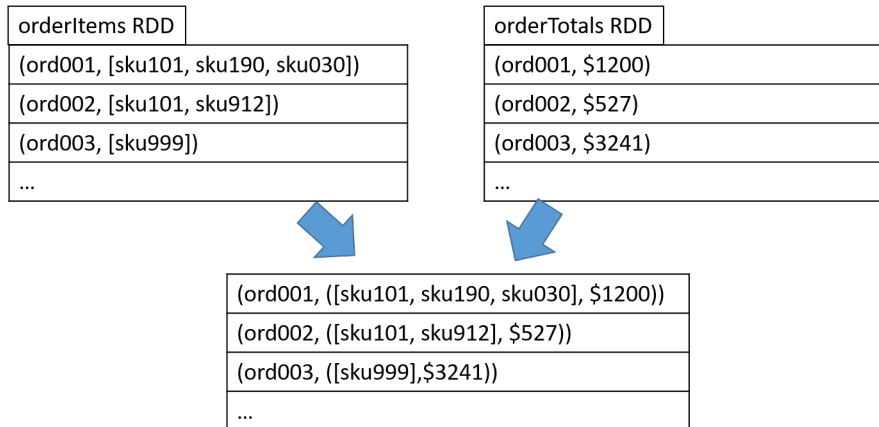
(ord003, sku999)
(ord002, sku101)
(ord002, sku912)
(ord001, sku030)
(ord001, sku190)
(ord001, sku101)
...



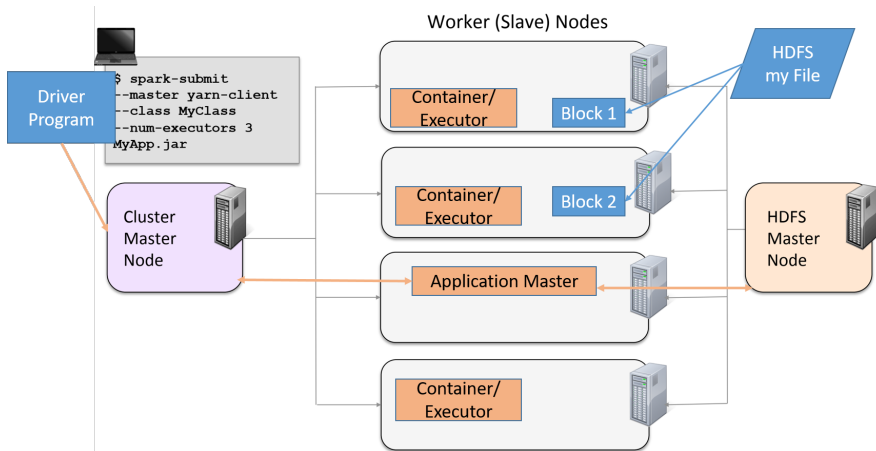
(ord001, [sku101, sku190, sku030,...])
(ord002, [sku101, sku912, ...])
(ord003, [sku999, ...])
...

## Example: join by key

Orders = orderItems.join(orderTotals)




# Running A Spark Job on YARN



# RDD Partitions

- Data is partitioned over the worker nodes
  - E.g., to follow blocks of an HDFS file
- Partitioning is done automatically by Spark
  - Optionally, you can control the number of partitions
  - You can specify the minimum number of partitions, default is 2
  - `sc.textFile("My File", 3)`



# Parallel Operations on Partitions

- Spark tries to maximize the localization of data processing
  - Group all transformations that can be processed on the same data partition
- Some transformations are partition-preserving
  - E.g., map, flatMap, filter
- Some transformations repartition
  - E.g., reduceByKey, groupByKey, sortByKey

- All operations that can work on the same data partition are grouped into a stage.
  - Tasks within a stage are pipelined together
- Spark divides the DAG of the job into stages
- How Spark Calculates Stages? Based on RDD dependencies
  - Narrow dependencies
    - Only one child depends on the RDD
    - No shuffle required
  - Wide (shuffle) dependencies
    - Multiple children depend on the RDD
    - Defines a new stage

\*DAG = Directed Acyclic Graph




## Example: Average Word Length By First Letter

We have the following chain of operations

```
Avglength = sc.textFile(file).flatMap(line: line.split())
```

```
.map(word:(word[0], len (word))).groupByKey()
```

```
.map((k , values) : (k, sum(values)/ len(values)))
```




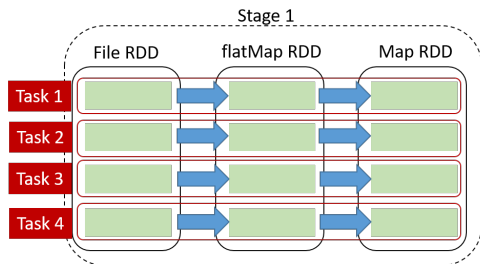
# Tasks Pipelining

We have the following chain of operations

```
Avglength = sc.textFile(file).flatMap(line: line.split())
```

```
.map(word:(word[0], len (word)).groupByKey()
```

```
.map((k , values) : (k, sum(values )/ len(values)))
```



# RDD Persistence

- Spark maintains lineage of RDDs by storing a reference to the parent RDD in the child one
- Each time an action is called on an RDD, Spark recursively traverses the lineage and performs the transformation
- This might be costly, especially in case of disk access
- Persistence makes Spark maintain the content of RDDs, default in memory
- Useful for iterative, e.g. machine learning, and interactive processing

